

Harshavardhan Raja

Senior Data Engineer | Cloud & Data Platforms | Gen AI-Augmented Development | 6+ Years Experience

Phone: +91-8008823497 | Email: harshaitg97@gmail.com | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY:

Results-oriented Senior Data Engineer with 6+ years of experience designing, building, and operating large-scale, production data platforms across AWS and GCP. Experienced practitioner of AI-assisted software development using Cursor AI — leveraging LLM-powered code generation, multi-file refactoring, and prompt engineering to accelerate engineering delivery. Strong expertise in data lakehouse architectures, Airflow-based orchestration, MongoDB, Snowflake, and S3-based data lakes. Demonstrated success owning complex data systems end-to-end and building reusable, configuration-driven frameworks that improve team productivity and data reliability.

CORE TECHNICAL SKILLS:

- **Languages:** Python, SQL
- **Cloud Platforms:** AWS (S3, MWAA, ECS, Lambda, CloudWatch, Secrets Manager), GCP (BigQuery, Dataflow, Pub/Sub, Cloud Composer, Cloud Run, Looker Studio)
- **Data Platforms & Databases:** MongoDB, Snowflake, BigQuery, MySQL, MS SQL Server
- **Data Engineering:** Apache Airflow, Airbyte, Apache Beam, Spark, ETL/ELT Pipelines, Incremental Sync, Data Validation, Schema Design
- **Gen AI & AI-Assisted Dev:** Cursor AI (codebase-aware code generation, multi-file refactoring, AI pair programming), Prompt Engineering, LLM-Augmented Development Workflows
- **Frameworks & APIs:** FastAPI, REST APIs, OAuth 2.0, JMESPath
- **DevOps & Tooling:** Docker, Terraform, Git, CI/CD (GitHub Actions)
- **Monitoring & Alerting:** Datadog, Cloud Monitoring

PROFESSIONAL EXPERIENCE:

Senior Data Engineer | Sequoia, India (Hybrid)

Sep 2024 - Apr 2026

- **AI-Augmented Engineering with Cursor AI:** Championed the adoption of Cursor AI as the team's primary development environment, establishing team-wide conventions for AI-assisted code generation, multi-file refactoring, and context-driven debugging across the Sequoia Cloud Data Exchange platform — reducing individual development cycle times by ~30%.
- **Prompt Engineering for Pipeline Generation:** Applied systematic prompt engineering within Cursor AI to auto-generate Airflow DAGs, API connector scaffolds, and data validation logic, enabling faster feature bootstrapping and establishing repeatable LLM-assisted coding patterns across the team.
- **AI-Accelerated ATS Integration Development:** Used Cursor AI's codebase-aware code generation to rapidly prototype and implement Workable and Lever ATS integrations — including JMESPath mapping logic, OAuth 2.0 handlers, and incremental sync state machines — cutting delivery time while maintaining code quality standards.
- **Multi-File AI Refactoring & Linting:** Leveraged Cursor AI's multi-file refactoring capabilities to systematically resolve large-scale linting and stability issues across the platform codebase, enabling high-velocity cleanup with zero regression impact and preserved test coverage.
- **LLM-Assisted Debugging & Incident Resolution:** Employed Cursor AI's @codebase context window to rapidly diagnose production issues in complex, multi-step pipelines — using AI-generated root cause hypotheses and targeted fix suggestions to reduce mean time to resolution (MTTR) during critical incidents.
- **AI Pair Programming for Framework Design:** Used Cursor AI as an active co-pilot to iteratively design the reusable MongoDB export framework, leveraging AI-suggested patterns for configuration-driven DAG templates, distributed locking, and idempotent re-run logic to compress design-to-implementation cycles.
- **Core Platform Contribution & ATS Incremental Sync:** Core contributor to the Sequoia Cloud Data Exchange / Lakehouse Platform ingesting and standardizing HR and ATS data from 20+ third-party systems for 50+ customer organizations. Designed ATS incremental sync system reducing data freshness from daily to near-hourly, with safe merging of incremental and historical data across candidates, applications, and jobs.
- **Reusable Export Framework & Platform Reliability:** Built a configuration-driven MongoDB → S3 → Snowflake export framework with dynamic Airflow orchestration, parallelism, and idempotent re-runs. Improved platform reliability via batch controls, distributed locking, and graceful failure handling for high-cardinality production datasets.
- **Cross-Functional Collaboration:** Collaborated with platform, analytics, and ops teams on ingestion standards, data models, and monitoring practices. Documented frameworks for faster onboarding. Recognized with the **Sequoia Backpack Award (2025)** for impactful contributions to enterprise data platform initiatives.

Senior Associate - Cloud Data Engineer | BlackBuck Insights, India

Dec 2021 - May 2024

- **Enterprise ETL Management:** Developed and maintained 50+ production ETL pipelines using Apache Airflow (Cloud Composer) for ingestion from Snowflake and GCS into BigQuery, supporting analytics and ML-driven use cases for a logistics and parcel shipping client.
- **Performance Optimization:** Improved pipeline performance and reliability, reducing average execution runtimes by ~30% through query optimization and workflow refactoring.
- **ML & Analytics Integration:** Integrated Vertex AI models into Airflow DAGs to enable automated demand forecasting and advanced analytics directly on data warehouse layers.
- **Microservices & Containerization:** Developed and operated 10+ FastAPI microservices exposing ML outputs to downstream applications with improved latency. Built and maintained Docker-based ML deployments on Cloud Run, reducing model release cycles by ~40%.
- **SLA Compliance & Monitoring:** Implemented Datadog-based monitoring and alerting across critical pipelines and services, successfully supporting 99.9% SLA compliance.
- **Tenant Onboarding Framework:** Created reusable DAG templates and onboarding documentation, reducing setup and configuration time for new tenants by ~40%.

Associate Consultant - Cloud Data Engineer | Virtusa, Hyderabad, India

Jul 2019 - Dec 2021

- **Enterprise Scale Pipelines:** Supported ~1,000 GCP-based production data pipelines as part of Google's External Workforce for a large mining client, ensuring continuous operation and high availability.
- **Large-Scale Cloud Migration:** Played a key role in migrating ~100 TB of legacy data from on-premise systems to BigQuery using Apache Beam and Cloud Dataflow.
- **Disaster Recovery Automation:** Built an automated BigQuery disaster recovery solution for schema and data backups, reducing recovery time objectives (RTO) from hours to minutes.
- **Infrastructure as Code (IaC):** Authored reusable Terraform modules to provision GCS, Compute Engine, and BigQuery resources aligned with CIS compliance and security standards.
- **Dashboards & ML R&D:** Built Looker Studio dashboards for platform health and KPI visibility for senior stakeholders. Collaborated with Virtusa X-Labs on a CNN-based pneumonia detection prototype.

EDUCATION & CERTIFICATIONS:

Education: B.Tech - Chemical Science & Technology, Indian Institute of Technology, Guwahati (Class of 2019)

Certifications:

- Google Cloud Certified - Professional Data Engineer
- Google Cloud Certified - Professional Cloud Architect

AWARDS & RECOGNITION:

- **Sequoia Backpack Award (2025)** – Recognized for impactful contributions to enterprise data platform initiatives.
- **BBi Core Value Award (Q3 2023)** – Awarded for outstanding collaboration and engineering reliability.
- **Virtusa Top Talent (H2 FY21) & Top Performer Bonus (2020-21)** – Awarded for excellence in technical delivery.
- **Virtusa V-Energiser Award (May 2021)** – Recognized for high performance and energy within the team.