

Harshavardhan Raja

Senior Data Engineer | Cloud & Data Platforms | 6+ Years Experience

Phone: +91-8008823497 | [Email: harshaitg97@gmail.com](mailto:harshaitg97@gmail.com) | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY:

Results-driven Senior Data Engineer with 6+ years of experience building, scaling, and operating production-grade data platforms across AWS and GCP. Proven track record of collaborating remotely with US-based and global organizations to architect high-performance data lakehouses, serverless workflows, and robust ETL/ELT pipelines. Specialized in building reusable, configuration-driven automation frameworks that boost developer productivity and data reliability. A self-driven engineer accustomed to working asynchronously across distributed international teams.

CORE TECHNICAL SKILLS:

- **Languages:** Python, SQL
- **Cloud Platforms:** AWS (S3, MWAA, ECS, Lambda, CloudWatch, Secrets Manager), GCP (BigQuery, Dataflow, Pub/Sub, Cloud Composer, Cloud Run, Looker Studio)
- **Data Platforms & Databases:** MongoDB, Snowflake, BigQuery, MySQL, MS SQL Server
- **Data Engineering:** Apache Airflow, Airbyte, Apache Beam, Spark, ETL/ELT Pipelines, Incremental Sync, Data Validation, Schema Design
- **Frameworks & APIs:** FastAPI, REST APIs, OAuth 2.0, JMESPath
- **DevOps & Tooling:** Docker, Terraform, Git, CI/CD (GitHub Actions, workflows), Cursor AI (AI-assisted development)
- **Monitoring & Alerting:** Datadog, Cloud Monitoring

PROFESSIONAL EXPERIENCE:

Senior Data Engineer | Sequoia (Hybrid)

Sep 2024 - Apr 2026

- **Core Platform Contribution:** Core contributor to the Sequoia Cloud Data Exchange / Lakehouse Platform, a multi-tenant data platform ingesting and standardizing HR and ATS data from 20+ third-party systems for 50+ customer organizations.
- **ATS Incremental Sync System:** Designed and implemented an ATS incremental sync system by enhancing the Generic Airbyte Sync framework to support hourly syncs, time-based incremental extraction, state lookback handling, and S3 partitioning. Successfully reduced data freshness from daily to near-hourly while safely merging incremental and historical data across candidates, applications, and jobs.
- **API & Integration Leadership:** Led end-to-end ATS integrations, including Workable ATS (custom API extraction, JMESPath mappings, CDM generation) and Lever ATS migration from legacy ingestion to Airbyte, implementing OAuth-based authentication, incremental syncs, and batch-optimized processing.
- **Reusable Export Framework:** Designed and built a reusable, configuration-driven MongoDB export framework to ingest curated platform data into AWS S3 and Snowflake, paired with dynamic Apache Airflow orchestration for configuration-driven execution, parallelism, idempotent re-runs, and automated downstream DAG triggering.
- **Platform Scale & Reliability:** Improved platform scalability and reliability by introducing batch controls, volume thresholds, distributed locking, and graceful failure handling, ensuring stable processing of large, high-cardinality datasets in production.
- **Quality & Stability Audits:** Drove platform reliability and code quality improvements, refactoring complex processing logic, resolving large-scale linting and stability issues, and stabilizing production pipelines with zero regression impact.
- **Distributed Collaboration:** Collaborated closely with platform, analytics, and operations teams to align on ingestion standards, data models, monitoring practices, and production readiness, while documenting frameworks to enable faster onboarding.

Freelance Data Engineering Consultant | Bonsai Data Solutions, LLC (US Remote)

Jan 2023 - Aug 2024 (Contract)

- **Serverless Engineering:** Engineered custom, automated data pipelines and serverless microservices using Google Cloud Functions and Cloud Run, optimizing execution latency and minimizing compute costs.
- **ETL & Reverse ETL Automation:** Developed, tested, and automated high-impact ETL and Reverse ETL data pipelines, orchestrating seamless data flows between client operational databases and cloud data warehouses.
- **Python App Testing:** Designed and executed robust automated testing frameworks for custom Python packages and application suites, elevating code coverage and streamlining deployment reliability.
- **BigQuery Data Warehousing:** Architected and optimized scalable BigQuery data warehouses for international clients, designing performant schemas and fine-tuning complex SQL queries to drive rapid analytics insights.

Senior Associate - Cloud Data Engineer | BlackBuck Insights

Dec 2021 - May 2024 (Remote)

- **US Client Engagement:** Served as a lead Data Engineer for a major US-based logistics and parcel shipping client, supporting high-throughput analytics and ML-driven use cases.
- **Enterprise ETL Management:** Developed and maintained 50+ production ETL pipelines using Apache Airflow (Cloud Composer) for ingestion from Snowflake and GCS into BigQuery, reducing average execution runtimes by ~30%.
- **ML & Analytics Integration:** Integrated Vertex AI models into Airflow DAGs to enable automated demand forecasting and advanced analytics directly on data warehouse layers.
- **Microservices & Containerization:** Developed and operated 10+ FastAPI microservices exposing ML outputs to downstream applications with improved latency. Built and maintained Docker-based ML deployments on Cloud Run, reducing model release cycles by ~40%.
- **SLA Compliance & Monitoring:** Implemented Datadog-based monitoring and alerting across critical pipelines and services, successfully supporting 99.9% SLA compliance.
- **Tenant Onboarding Framework:** Created reusable DAG templates and onboarding documentation, reducing setup and configuration time for new tenants by ~40%.

Associate Consultant - Cloud Data Engineer | Virtusa, Hyderabad

Jul 2019 - Dec 2021

- **Google External Workforce:** Supported ~1,000 GCP-based production data pipelines as part of Google's External Workforce for a large mining client, maintaining high availability and uptime.
- **Large-Scale Cloud Migration:** Played a key role in migrating ~100 TB of legacy data from on-premise systems to BigQuery using Apache Beam and Cloud Dataflow.
- **Disaster Recovery Automation:** Built an automated BigQuery disaster recovery solution for schema and data backups, reducing recovery time objectives (RTO) from hours to minutes.
- **IoT Real-Time Ingestion:** Designed and implemented IoT data pipelines for real-time sensor ingestion and analytics, integrated with Virtusa's Cognisense platform.
- **Infrastructure as Code (IaC):** Authored reusable Terraform modules to provision GCS, Compute Engine, and BigQuery resources aligned with CIS compliance and security standards.
- **Dashboards & ML R&D:** Built Looker Studio dashboards to provide clear visibility into data platform health and KPIs for senior stakeholders. Collaborated with Virtusa X-Labs on applying ML to healthcare analytics, contributing to a CNN-based pneumonia detection prototype.

EDUCATION & CERTIFICATIONS:

- **Education:** B.Tech - Chemical Science & Technology, Indian Institute of Technology, Guwahati (Class of 2019)
- **Certifications:**
 - Google Cloud Certified - Professional Data Engineer
 - Google Cloud Certified - Professional Cloud Architect

AWARDS & RECOGNITION:

- **Sequoia Backpack Award (2025)** - Recognized for impactful contributions to enterprise data platform initiatives.
- **BBI Core Value Award (Q3 2023)** - Awarded for outstanding collaboration and engineering reliability.
- **Virtusa Top Talent (H2 FY21) & Top Performer Bonus (2020-21)** - Awarded for excellence in technical delivery.
- **Virtusa V-Energiser Award (May 2021)** - Recognized for high performance and energy within the team.